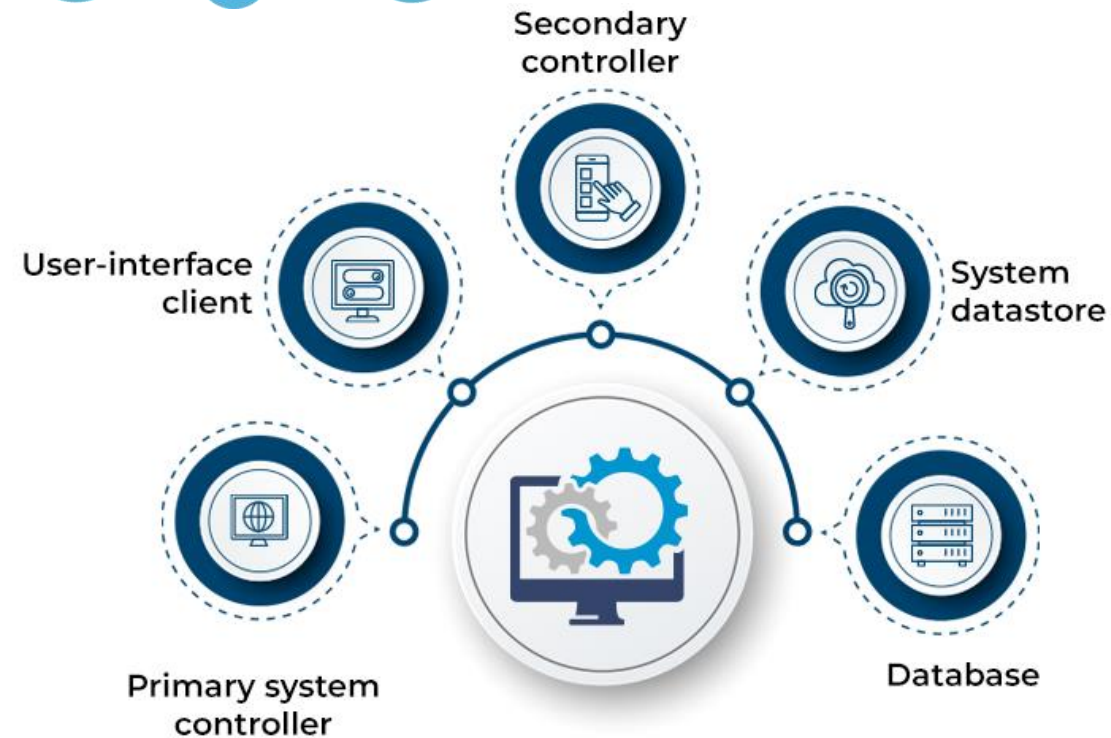




Sistem Terdistribusi

IF2222

04: RPC

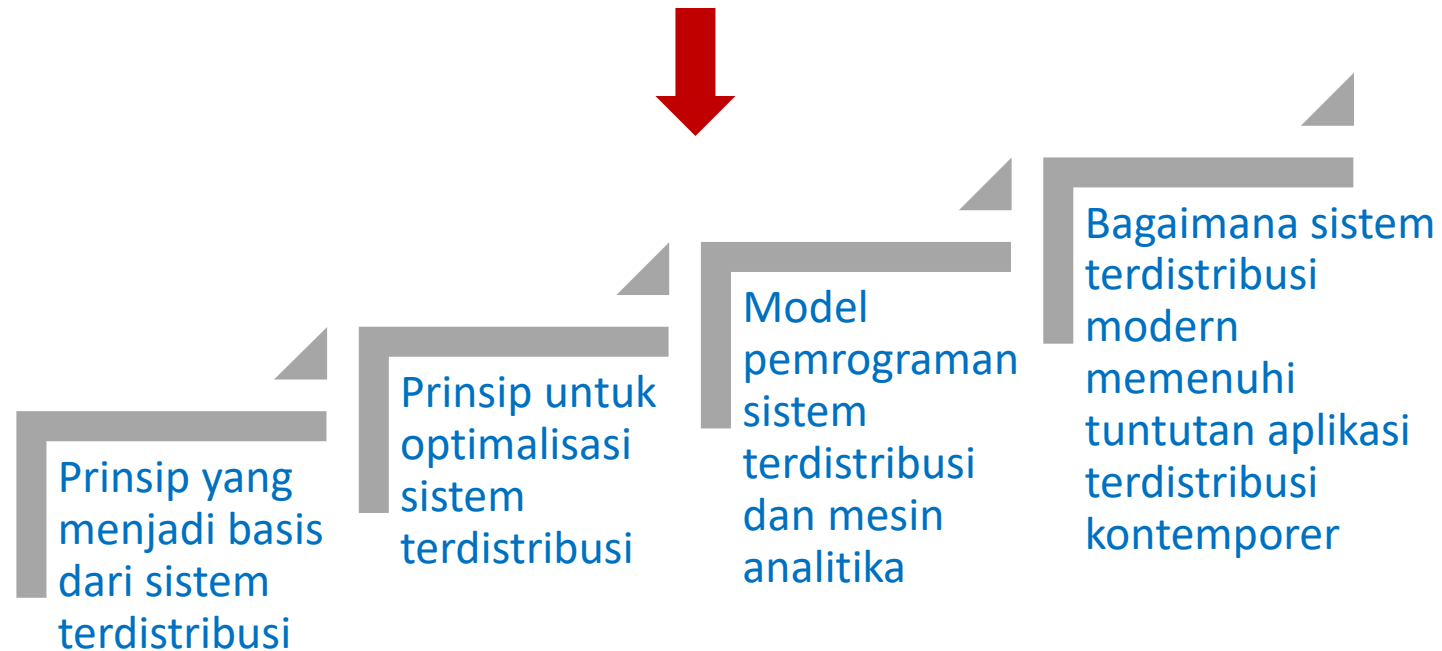


Sistem Terdistribusi 2022

1. Mengenal Sistem Terdistribusi
2. Review Jaringan Komputer (layer 2, 3, dan 4)
3. Arsitektur Sistem Terdistribusi
- 4. *Remote Procedure Calls (RPC)***
5. Layanan Penamaan
6. Sinkronisasi Data (2 pekan)
7. *Message Passing Interface (MPI)*
8. Contoh Arsitektur: Hadoop, Pregel, Blockchain
9. Teknik *Caching*
10. Teknik Replikasi Data (2 pekan)
11. Basis Data Terdistribusi
12. Toleransi Kegagalan

Capaian Pembelajaran

Kuliah ini bertujuan memberikan pemahaman mendalam dan pengalaman langsung tentang:



Today...

- Last Session:

- Arsitektur Sistem Terdistribusi

- Today's Session:

- Remote Procedure Calls- Part I
 - Sockets
 - Remote Invocations

- Announcement:

Communicating Entities in Distributed Systems

- Communicating entities in distributed systems can be classified into two types:
 - System-oriented entities
 - Processes
 - Threads
 - Nodes
 - Problem-oriented entities
 - Objects (in *object-oriented programming* based approaches)

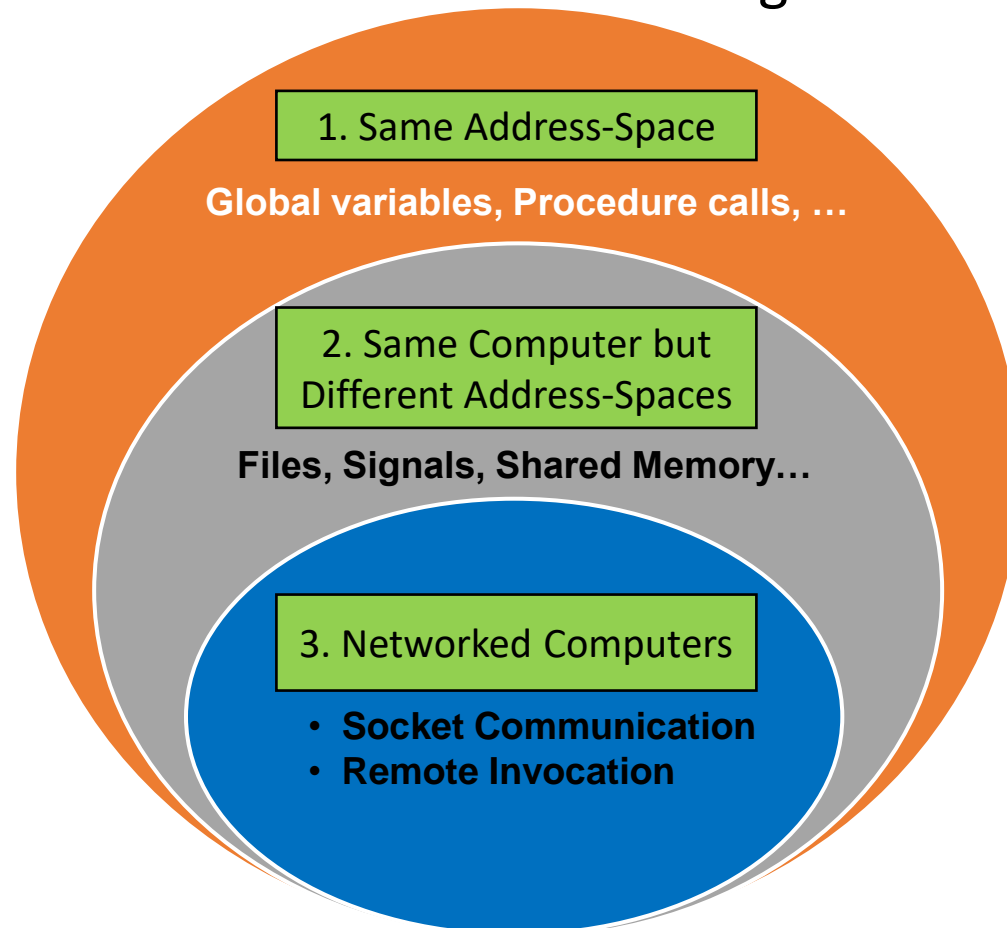
How can entities in distributed systems communicate?

Communication Paradigms

- Communication paradigms describe and classify a set of methods by which entities can interact and exchange data

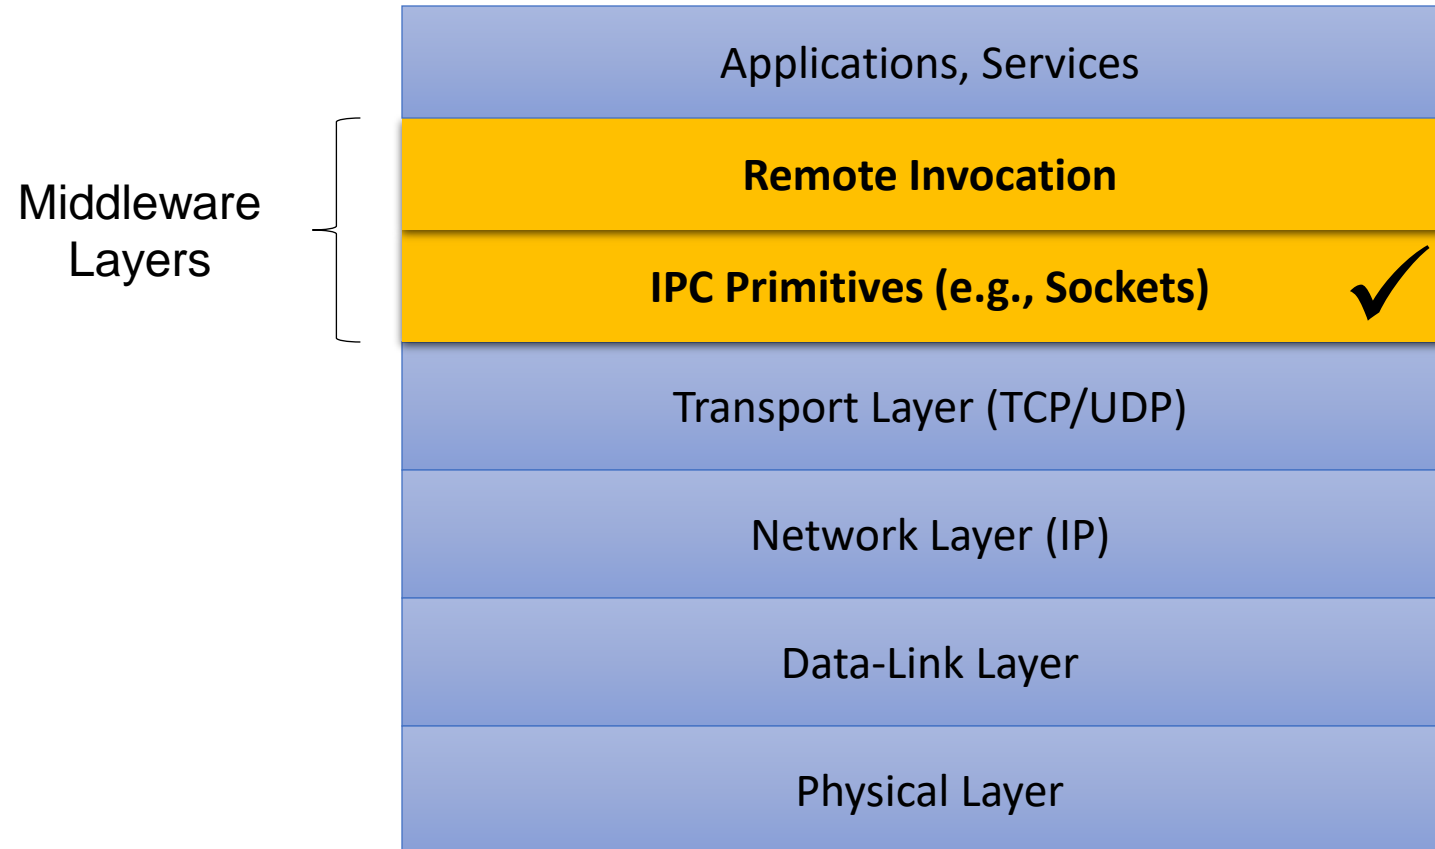
Classification of Communication Paradigms

- Communication paradigms can be categorized into *three* types based on where the entities reside. If entities are running on:



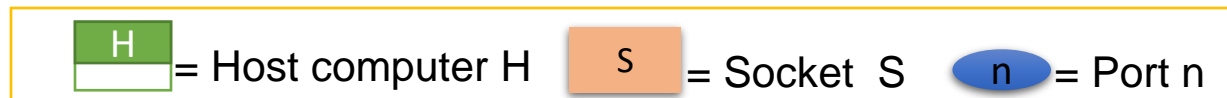
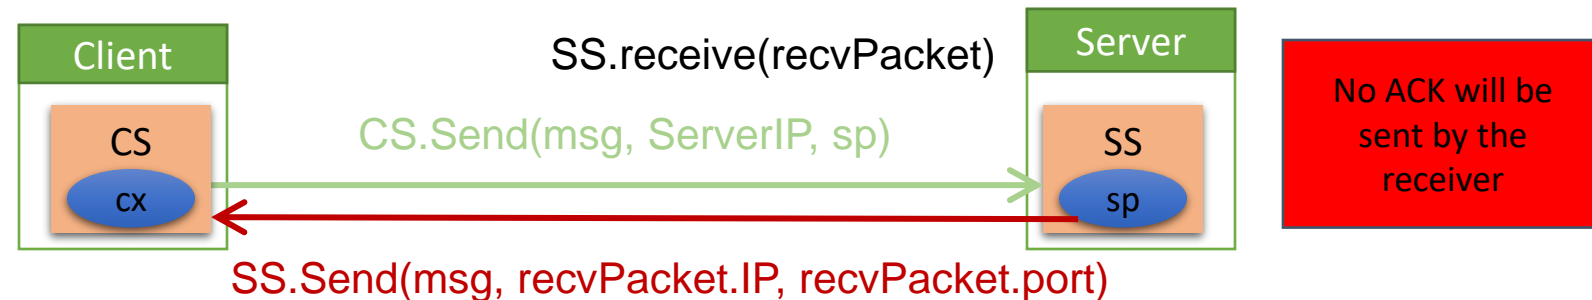
Today, we will study how entities that reside on **networked computers** communicate in distributed systems using socket communication and remote invocation

Middleware Layers



UDP Sockets

- UDP provides *connectionless* communication, with no acknowledgements or message retransmissions
- Communication mechanism:
 - Server opens a UDP socket *SS* at a known port *sp*,
 - Socket *SS* waits to receive a request
 - Client opens a UDP socket *CS* at a random port *cx*
 - Client socket *CS* sends a message to *ServerIP* and port *sp*
 - Server socket *SS* may send back data to *CS*

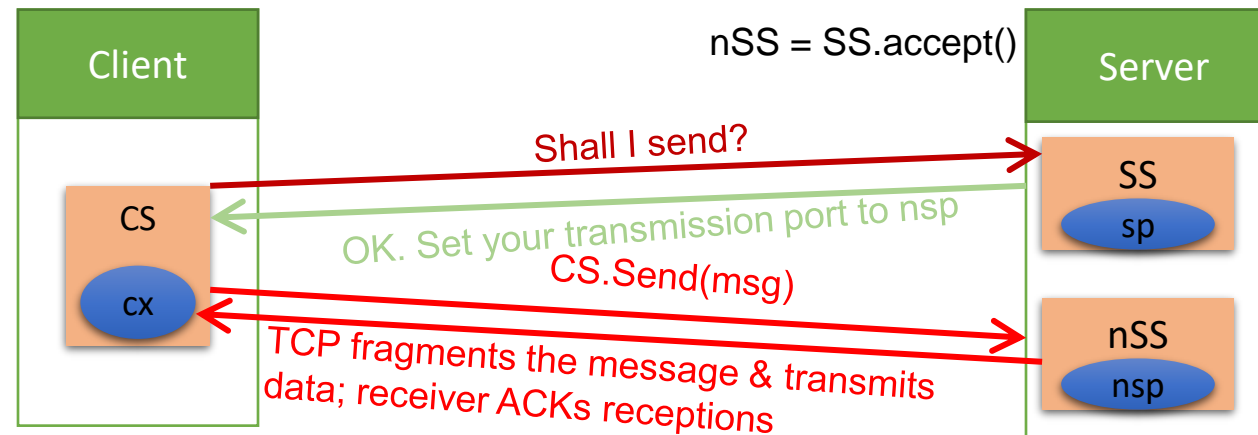


UDP– Design Considerations

- Sender must explicitly fragment a long message into smaller chunks before transmission
 - A maximum size of 548 bytes is suggested for transmission
- Messages may be delivered out-of-order
 - If necessary, programmer must re-order packets
- Communication is not reliable
 - Messages might be dropped due to check-sum errors or buffer overflows at routers
- Receiver should allocate a buffer that is big enough to fit the sender's message
 - Otherwise the message will be truncated

TCP Sockets

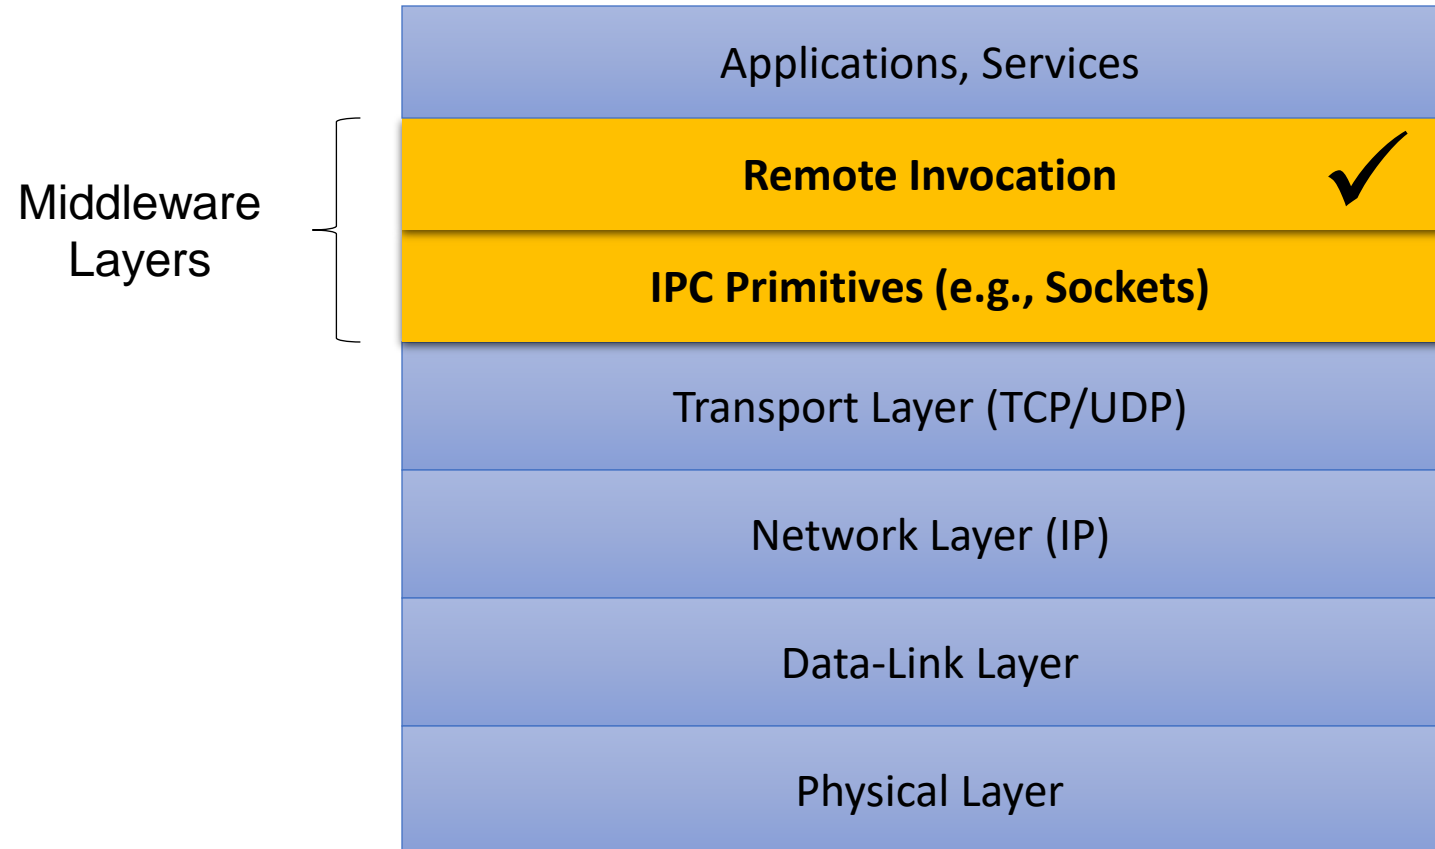
- TCP provides *in-order* delivery, *reliability*, and *congestion control*
- Communication mechanism:
 - Server opens a TCP server socket *SS* at a known port *sp*
 - Server waits to receive a request (using *accept* call)
 - Client opens a TCP socket *CS* at a random port *cx*
 - *CS* initiates a **connection initiation message** to ServerIP and port *sp*
 - Server socket *SS* allocates a **new socket** *NSS* on **random port** *nsp* for the client
 - *CS* can **send data** to *NSS*



Main Advantages of TCP

- TCP ensures in-order delivery of messages
- Applications can send messages of any size
- TCP ensures *reliable communication* via using acknowledgements and retransmissions
- Congestion control of TCP regulates sender rate, and thus prevents network overload

Middleware Layers

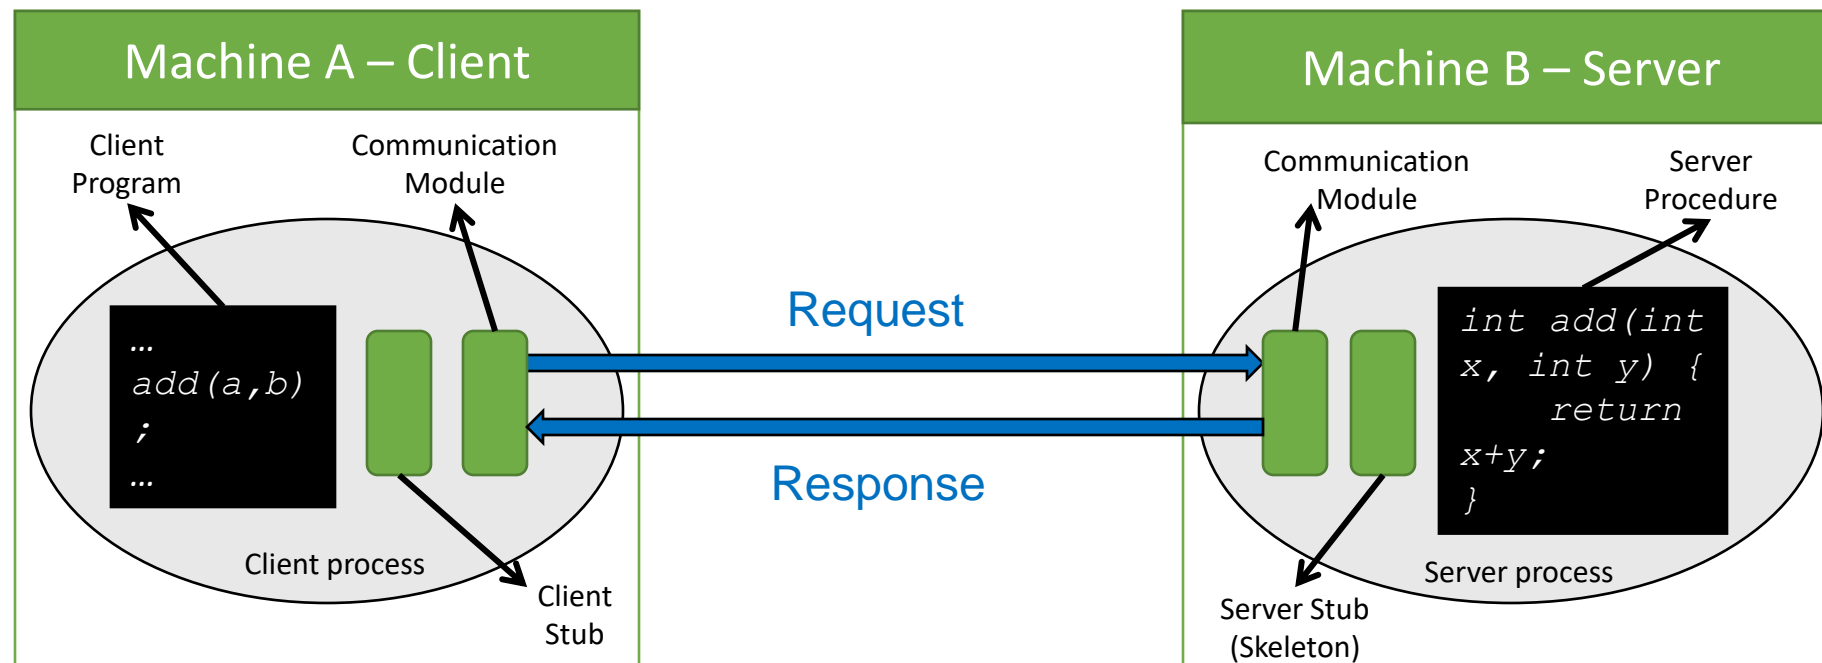


Remote Invocation

- Remote invocation enables an entity to call a procedure that typically executes on another computer **without the programmer explicitly coding the details of communication**
 - The underlying middleware will take care of raw-communication
 - Programmer can *transparently* communicate with remote entity
- We will study two types of remote invocations:
 - a. Remote Procedure Calls (RPC)
 - b. Remote Method Invocation (RMI)

Remote Procedure Calls (RPC)

- RPC enables a sender to communicate with a receiver using a simple procedure call
 - No communication or message-passing is visible to the programmer
- Basic RPC Approach:



Client Stub

- The client stub:
 - Gets invoked by user code as a local procedure
 - *Packs* (or *serializes* or *marshals*) parameters into a request packet (say, request-pkt)
 - Invokes a client side transport routine (e.g., `makerpc(request-pkt, &reply-pkt)`)
 - *Unpacks* (or *de-serializes* or *unmarshals*) reply-pkt into output parameters
 - Returns to user code

Server Stub

- The server stub:
 - Gets invoked after a server side transport routine (e.g., `getrequest()`) is returned
 - Unmarshals arguments, de-multiplexes opcode, and invokes local server code
 - Marshals arguments, invokes a server-side transport routine (e.g., `sendresponse()`), and returns to server loop
 - E.g., Typical server main loop:

```
while (1) {  
    get-request (&p);    /* blocking call */  
    execute-request (p); /* demux based on opcode */  
}
```

Challenges in RPC

- Parameter passing via marshaling
 - Procedure parameters and results have to be transferred over the network as bits
- Data representation
 - Data representation has to be uniform
 - Architecture of the sender and receiver machines may differ
- Failure Independence
 - Client and server might fail independently

Challenges in RPC

- **Parameter passing via marshaling**
 - Procedure parameters and results have to be transferred over the network as bits
- Data representation
 - Data representation has to be uniform
 - Architecture of the sender and receiver machines may differ
- Failure Independence
 - Client and server might fail independently

Parameter Passing via Marshaling

- Packing parameters into a message that will be transmitted over the network is called *parameter marshalling*
- The parameters to the procedure and the result have to be marshaled before transmitting them over the network
- Two types of parameters can be passed:
 1. Value parameters
 2. Reference parameters

1. Passing Value Parameters

- Value parameters have complete information about the variable, and can be directly encoded into the message
 - E.g., integer, float, character
- Values are passed through call-by-value
 - The changes made by the callee procedure are not reflected in the caller procedure

2. Passing Reference Parameters

- Passing reference parameters like value parameters in RPC leads to incorrect results due to two reasons:
 - a. Invalidity of reference parameters at the server
 - Reference parameters are valid only within client's address space
 - **Solution:** Pass the reference parameter by copying the data that is referenced
 - b. Changes to reference parameters are not reflected back at the client
 - **Solution:** "Copy/Restore" the data
 - Copy the data that is referenced by the parameter
 - Copy-back the value at server to the client

Challenges in RPC

- Parameter passing via marshaling
 - Procedure parameters and results have to be transferred over the network as bits
- Data representation
 - Data representation has to be uniform
 - Architecture of the sender and receiver machines may differ
- Failure Independence
 - Client and server might fail independently

Data Representation

- Computers in DSs often have different architectures and operating systems
 - The size of the data-type differ
 - E.g., A *long* data-type is 4-bytes in 32-bit Unix, while it is 8-bytes in 64-bit Unix systems
 - The format in which the data is stored differs
 - E.g., Intel stores data in little-endian format, while SPARC stores in big-endian format
- The client and server have to agree on how simple data is represented in the message
 - E.g., Format and size of data-types such as integer, char and float

Challenges in RPC

- Parameter passing via marshaling
 - Procedure parameters and results have to be transferred over the network as bits
- Data representation
 - Data representation has to be uniform
 - Architecture of the sender and receiver machines may differ
- Failure Independence
 - Client and server might fail independently

Failure Independence

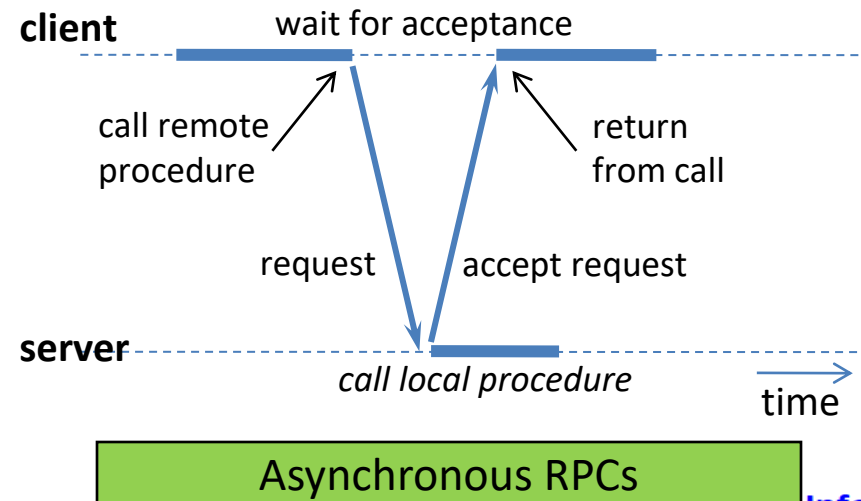
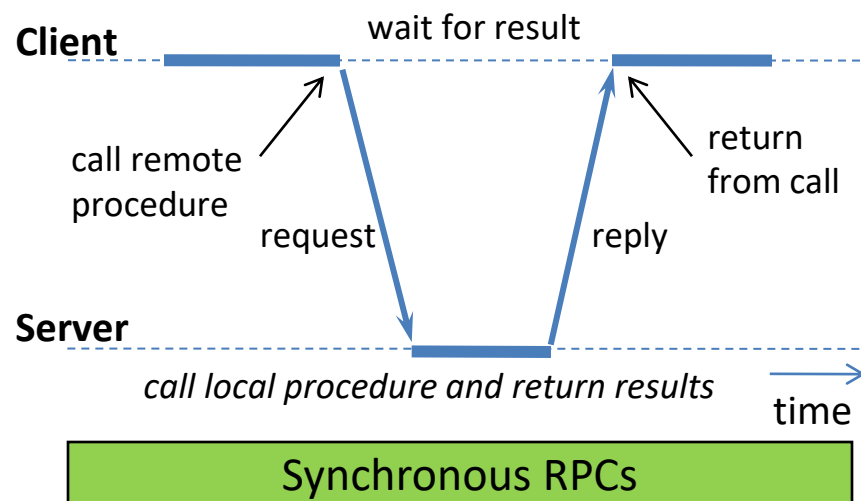
- In the local case, the client and server live or die together
- In the remote case, the client sees new *failure types* (more on this next lecture)
 - Network failure
 - Server machine crash
 - Server process crash
- Thus, failure handling code has to be more thorough (and essentially more complex)

Remote Procedure Call Types

- Remote procedure calls can be:
 - Synchronous
 - Asynchronous (or Deferred Synchronous)

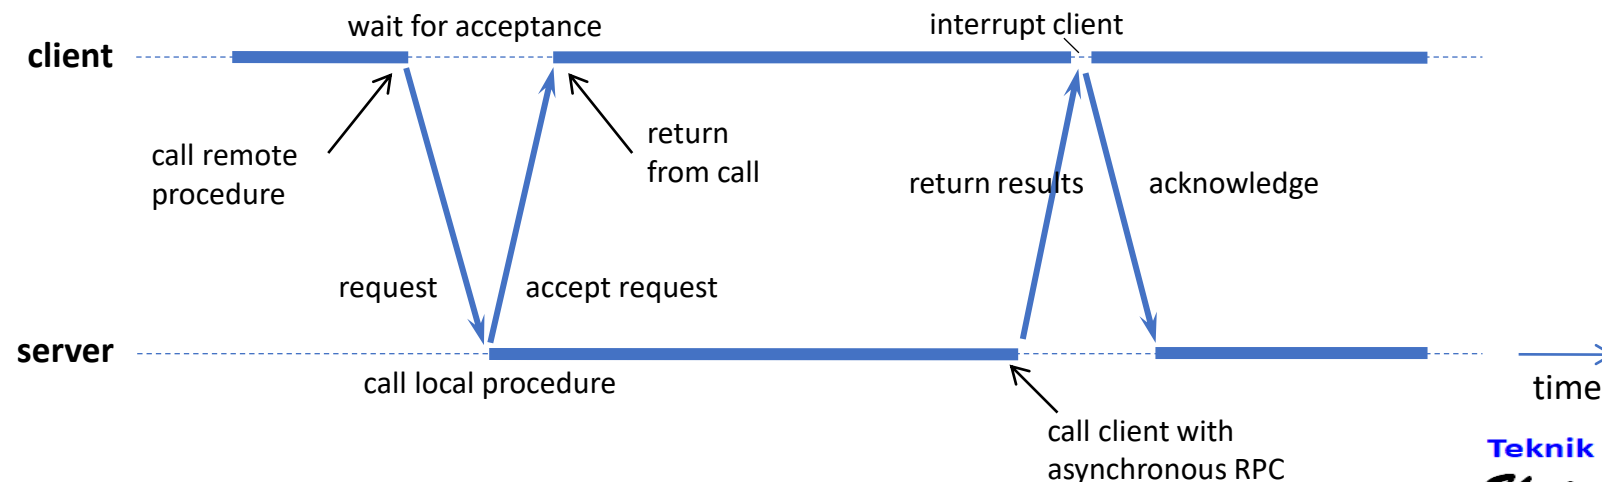
Synchronous vs. Asynchronous RPCs

- An RPC with strict request-reply blocks the client until the server returns
 - Blocking wastes resources at the client
- Asynchronous RPCs are used if the client does not need the result from server
 - The server immediately sends an ACK back to the client
 - The client continues the execution after an ACK from the server



Deferred Synchronous RPCs

- Asynchronous RPC is also useful when a client wants the results, but does not want to be blocked until the call finishes
- Client uses *deferred synchronous* RPCs
 - Single request-response RPC is split into two RPCs
 - First, client triggers an asynchronous RPC on server
 - Second, on completion, server calls-back client to deliver the results



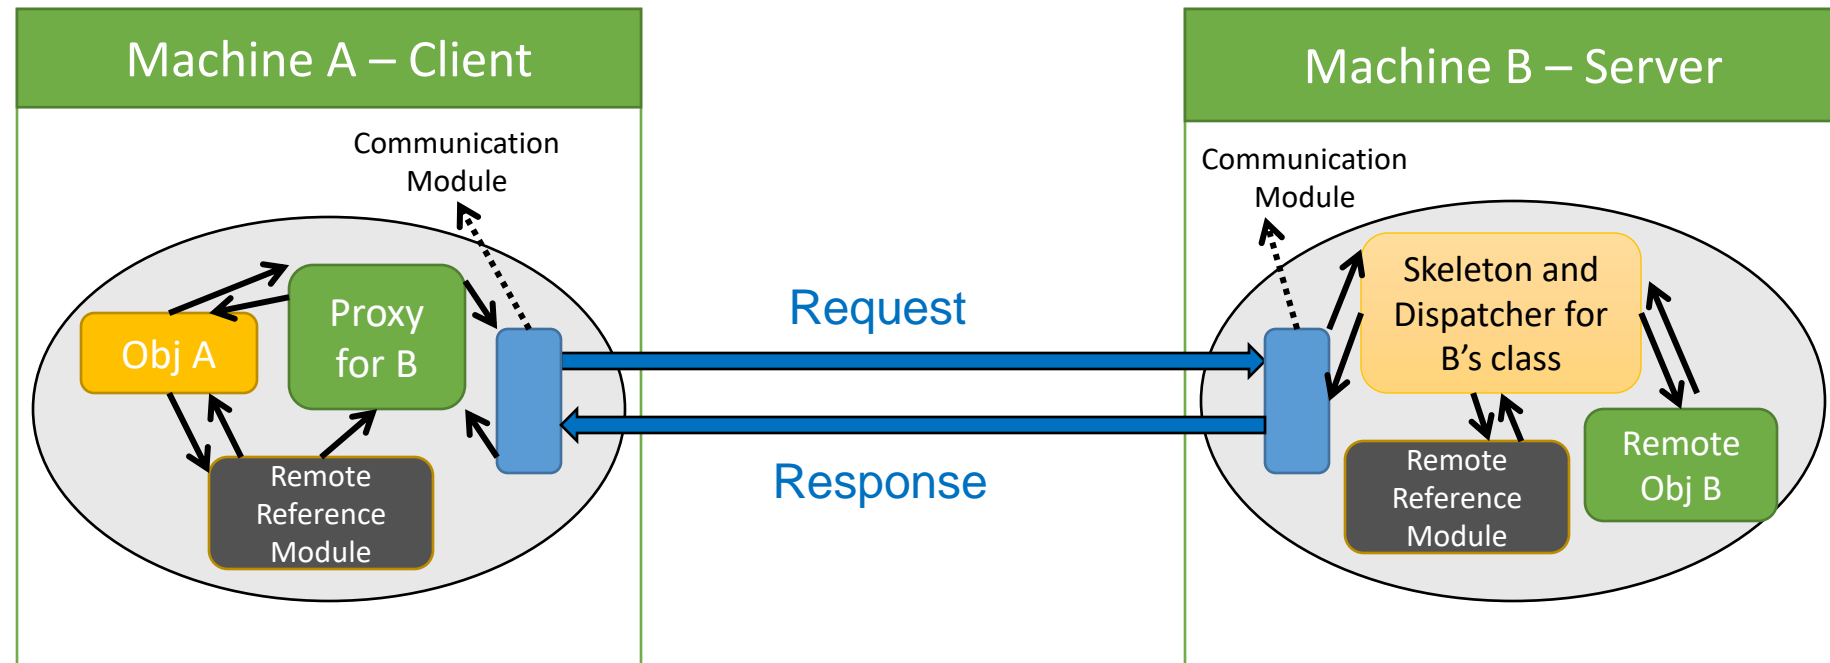
Remote Method Invocation (RMI)

- RMI is similar to RPC, but in a world of distributed objects
 - The programmer can use the full expressive power of object-oriented programming
 - RMI not only allows to pass value parameters, but also pass object references
- In RMI, a calling object can invoke a method on a potentially remote object

Remote Objects and Supporting Modules

- In RMI, objects whose methods can be invoked remotely are known as “*remote objects*”
 - Remote objects implement remote interfaces
- During any method call, the system has to resolve whether the method is being called on a local or a remote object
 - Local calls should be called on a local object
 - Remote calls should be called via remote method invocation
 - *Remote Reference Module* is responsible for translating between local and remote object references

RMI Control Flow



Today...

- Last Session:

- RPC- Part I

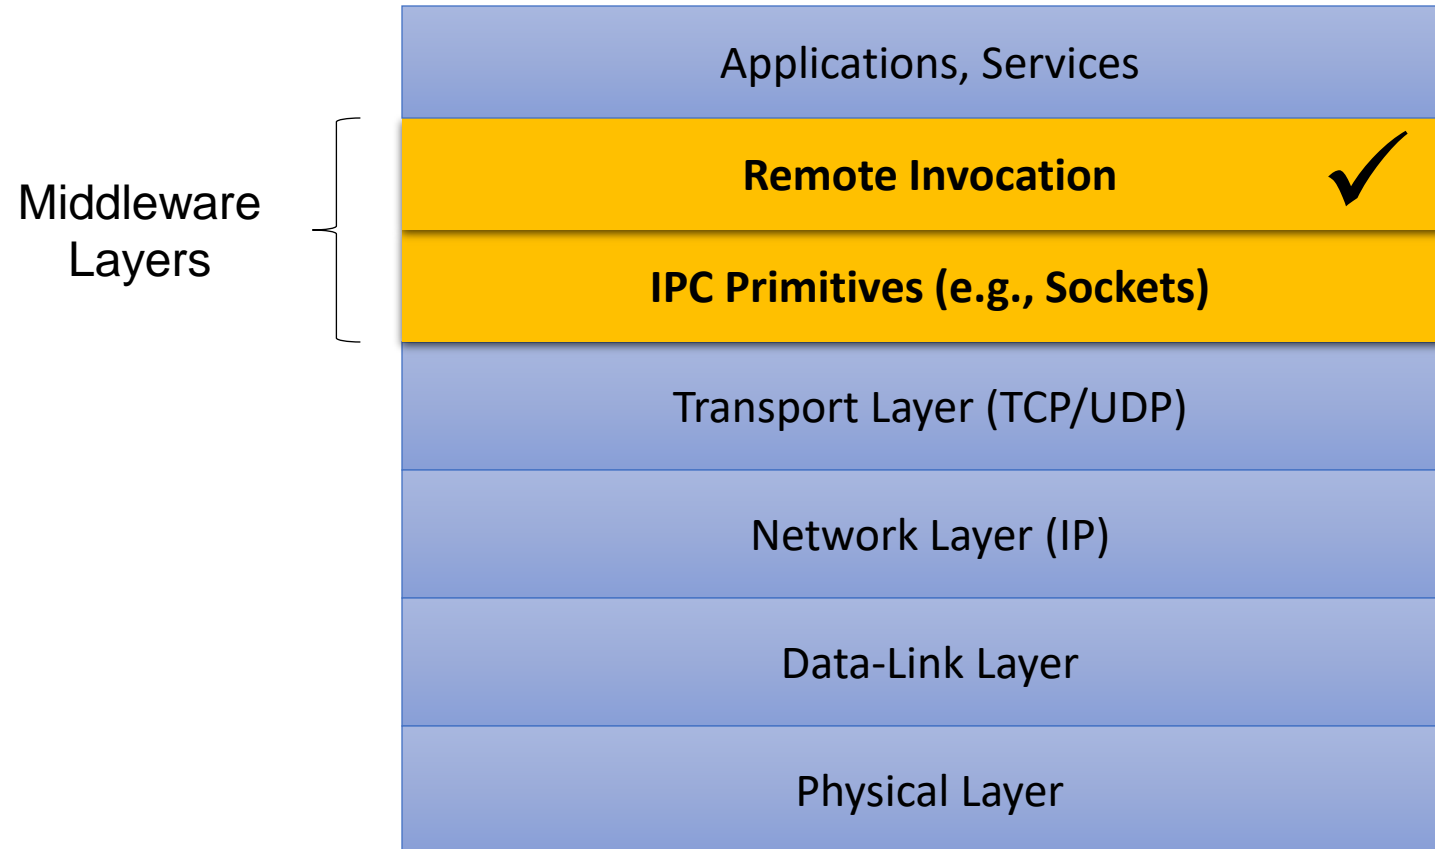
- Today's Session:

- Continue with Remote Procedure Calls

- Announcement:

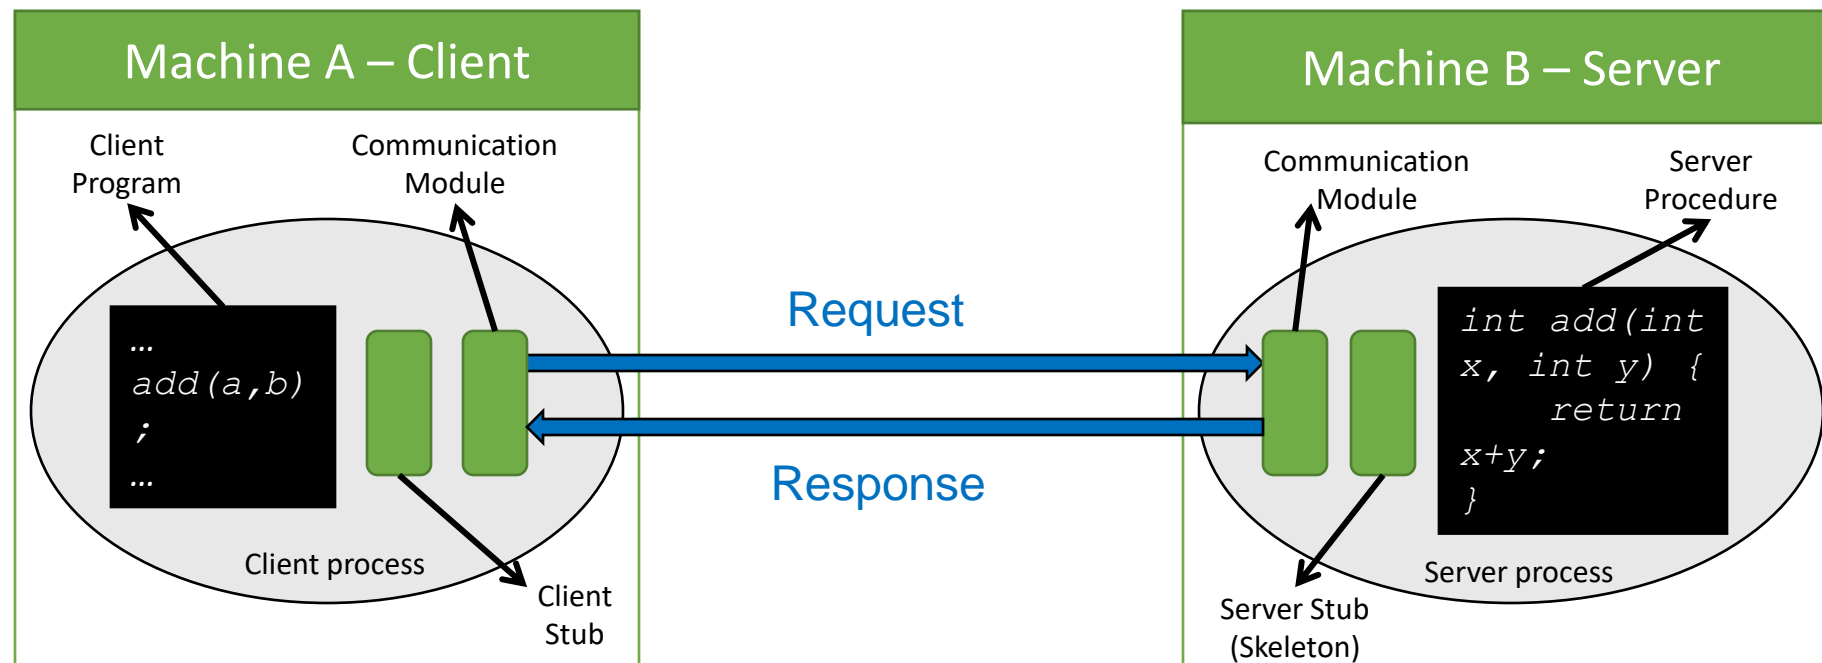
- Project I is now out. It is due on Feb 21 (design report is due on Feb 6)

Middleware Layers



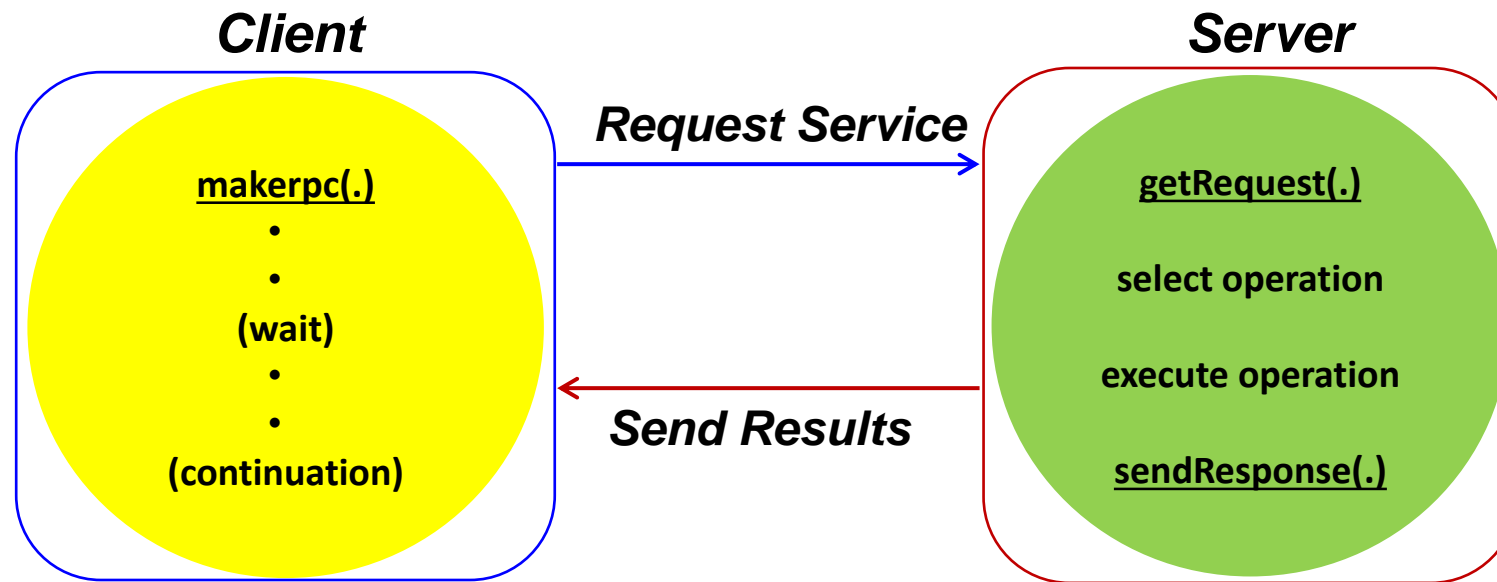
Remote Procedure Calls (RPC)

- RPC enables a sender to communicate with a receiver using a simple procedure call
 - No communication or message-passing is visible to the programmer
- Basic RPC Approach:



Transport Primitives

- RPC communication module (or *transport*) is mainly based on a trio of communication primitives, *makerpc(.)*, *getRequest(.)*, and *sendResponse(.)*



Failure Types

- RPC systems may suffer from various types of failures

Type of Failure	Description
<ul style="list-style-type: none">• Crash Failure	<ul style="list-style-type: none">• A server halts, but was working correctly until it stopped
<ul style="list-style-type: none">• Omission Failure<ul style="list-style-type: none">• Receive Omission• Send Omission	<ul style="list-style-type: none">• A server fails to respond to incoming requests<ul style="list-style-type: none">• A server fails to receive incoming messages• A server fails to send messages
<ul style="list-style-type: none">• Timing Failure	<ul style="list-style-type: none">• A server's response lies outside the specified time interval
<ul style="list-style-type: none">• Response Failure<ul style="list-style-type: none">• Value Failure• State Transition Failure	<ul style="list-style-type: none">• A server's response is incorrect<ul style="list-style-type: none">• The value of the response is wrong• The server deviates from the correct flow of control
<ul style="list-style-type: none">• Byzantine Failure	<ul style="list-style-type: none">• A server may produce arbitrary responses at arbitrary times

Timeout Mechanism

- To allow for occasions where a request or a reply message is lost, *makerpc(.)* can use a *timeout mechanism*
- There are various options as to what *makerpc(.)* can do after a timeout:
 - Either return immediately with an indication to the client that the request has failed
 - Or *retransmit* the request repeatedly until either a reply is received or the server is assumed to have failed
- How to pick a timeout value?
 - At best, use empirical/theoretical statistics
 - At worst, no good value exists

Idempotent Operations

- In cases when the request message is retransmitted, the server may receive it *more than once*
- This can cause an operation to be executed more than once for the same request
- Caveat: *Not* every operation can be executed more than once and obtain the same result each time!
- Operations that CAN be executed repeatedly with the same effect are called *idempotent operations*

Duplicate Filtering

- To avoid problems with operations, the server should:
 - Identify successive messages from the “same” client
 - Monotonically increasing *sequence numbers* can be used
 - Filter out duplicates
- Upon receiving a “duplicate” request, the server can:
 - Either **re-execute** the operation again and reply
 - Possible only for idempotent operations
 - Or **avoid re-executing** the operation via retaining its output in a non-volatile history (or *log*) file
 - Might necessitate *transactional semantics* (more on this later in the course)

Implementation Choices

- RPC transport can be implemented in different ways to provide different *delivery guarantees*. The main choices are:
 1. **Retry request service** (*client side*): Controls whether to retransmit the request service until either a reply is received or the server is assumed to have failed
 2. **Duplicate filtering** (*server side*): Controls when retransmissions are used and whether to filter out duplicate requests at the server
 3. **Retention of results** (*server side*): Controls whether to keep a history of result messages so as to enable lost replies to be retransmitted without re-executing the operations at the server

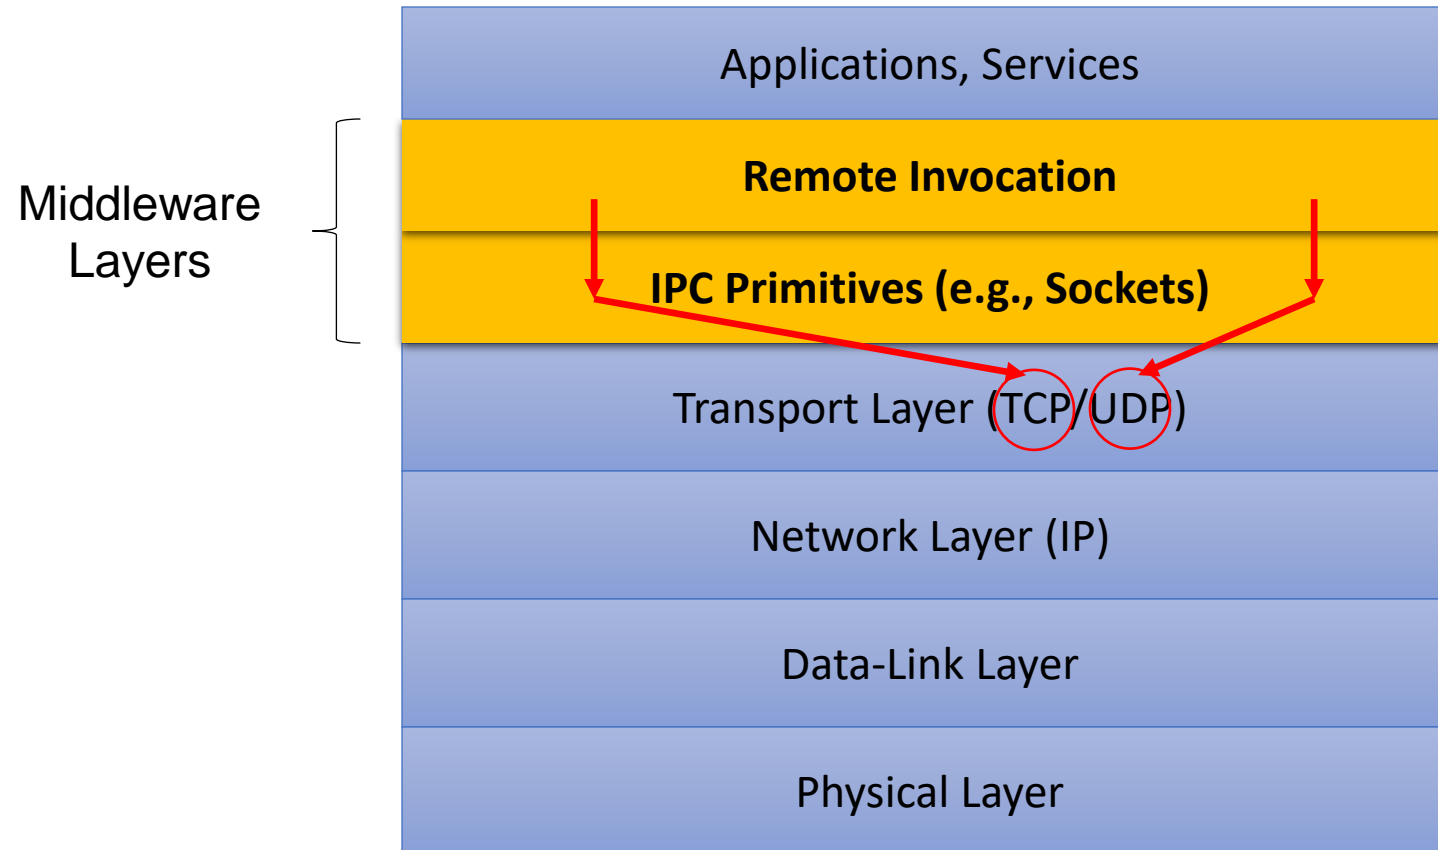
RPC Call Semantics

- Combinations of measures lead to a variety of possible *semantics* for the reliability of RPC

Fault Tolerance Measure			Call Semantics (Pertaining to Remote Procedures)
Retransmit Request Message	Duplicate Filtering	Re-execute Procedure or Retransmit Reply	
No	N/A	N/A	<i>Maybe</i>
Yes	No	Re-execute Procedure	<i>At-least-once</i>
Yes	Yes	Retransmit Reply	<i>At-most-once</i>

Ideally, we would want an *exactly-once* semantic!

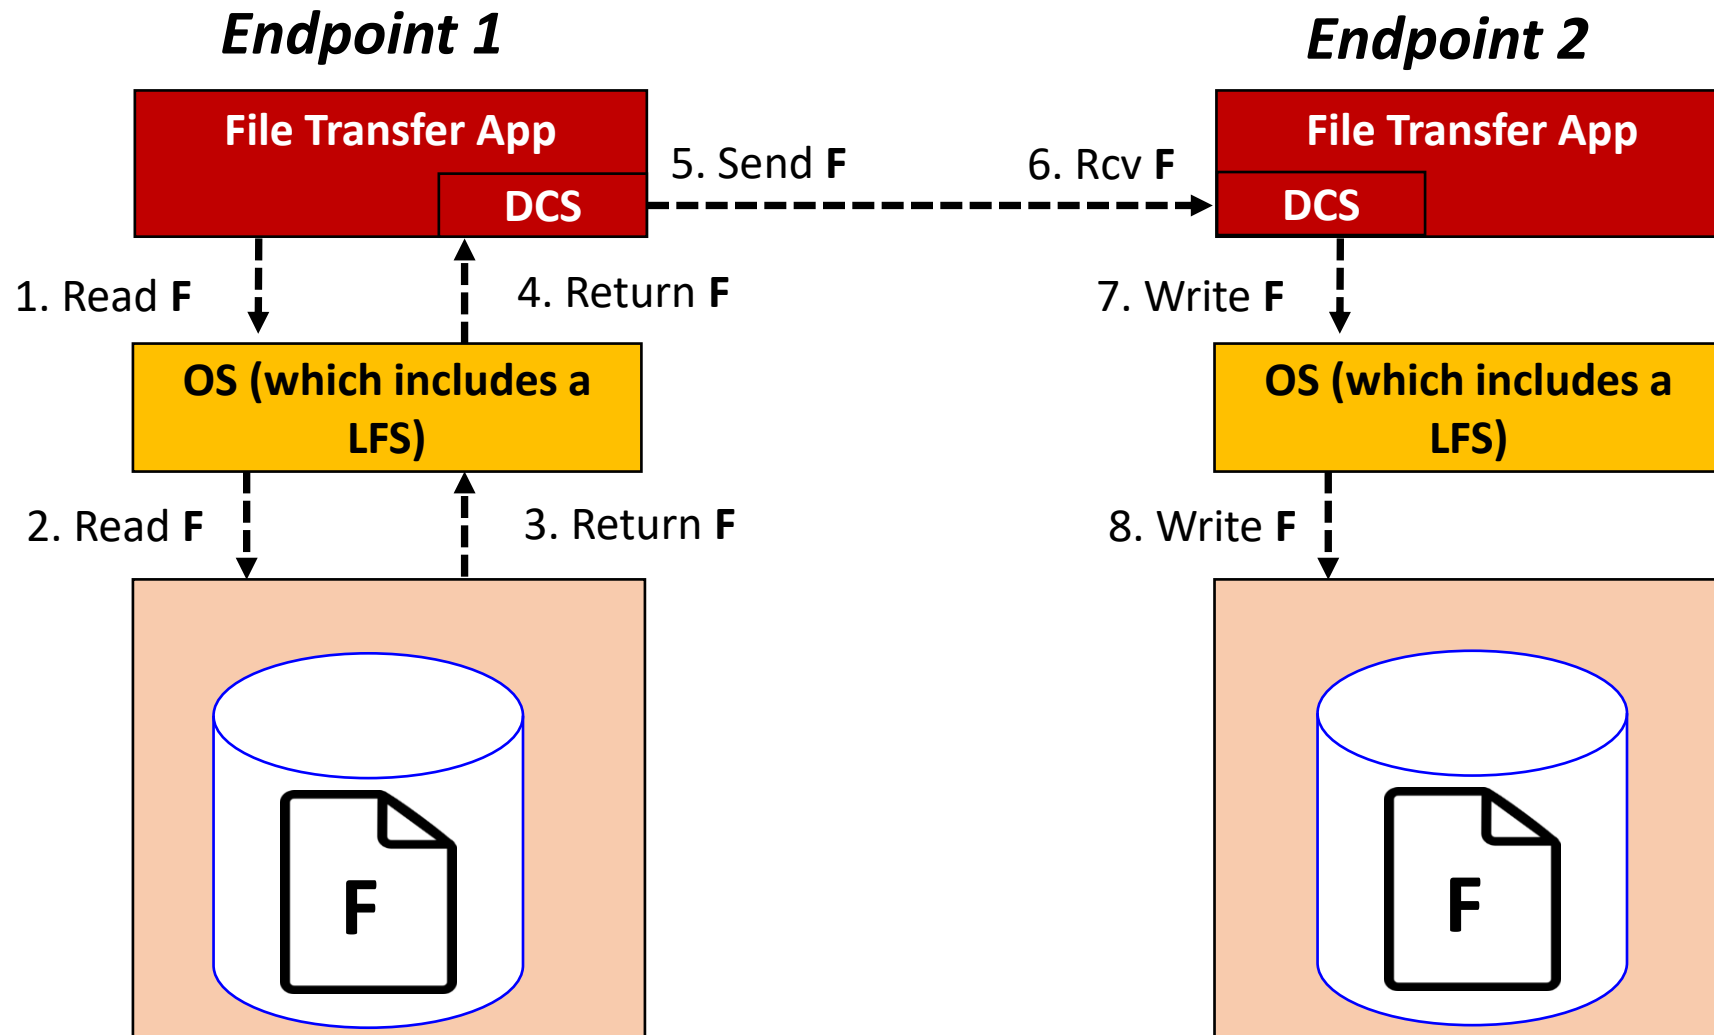
Middleware Layers



RPC over UDP or TCP

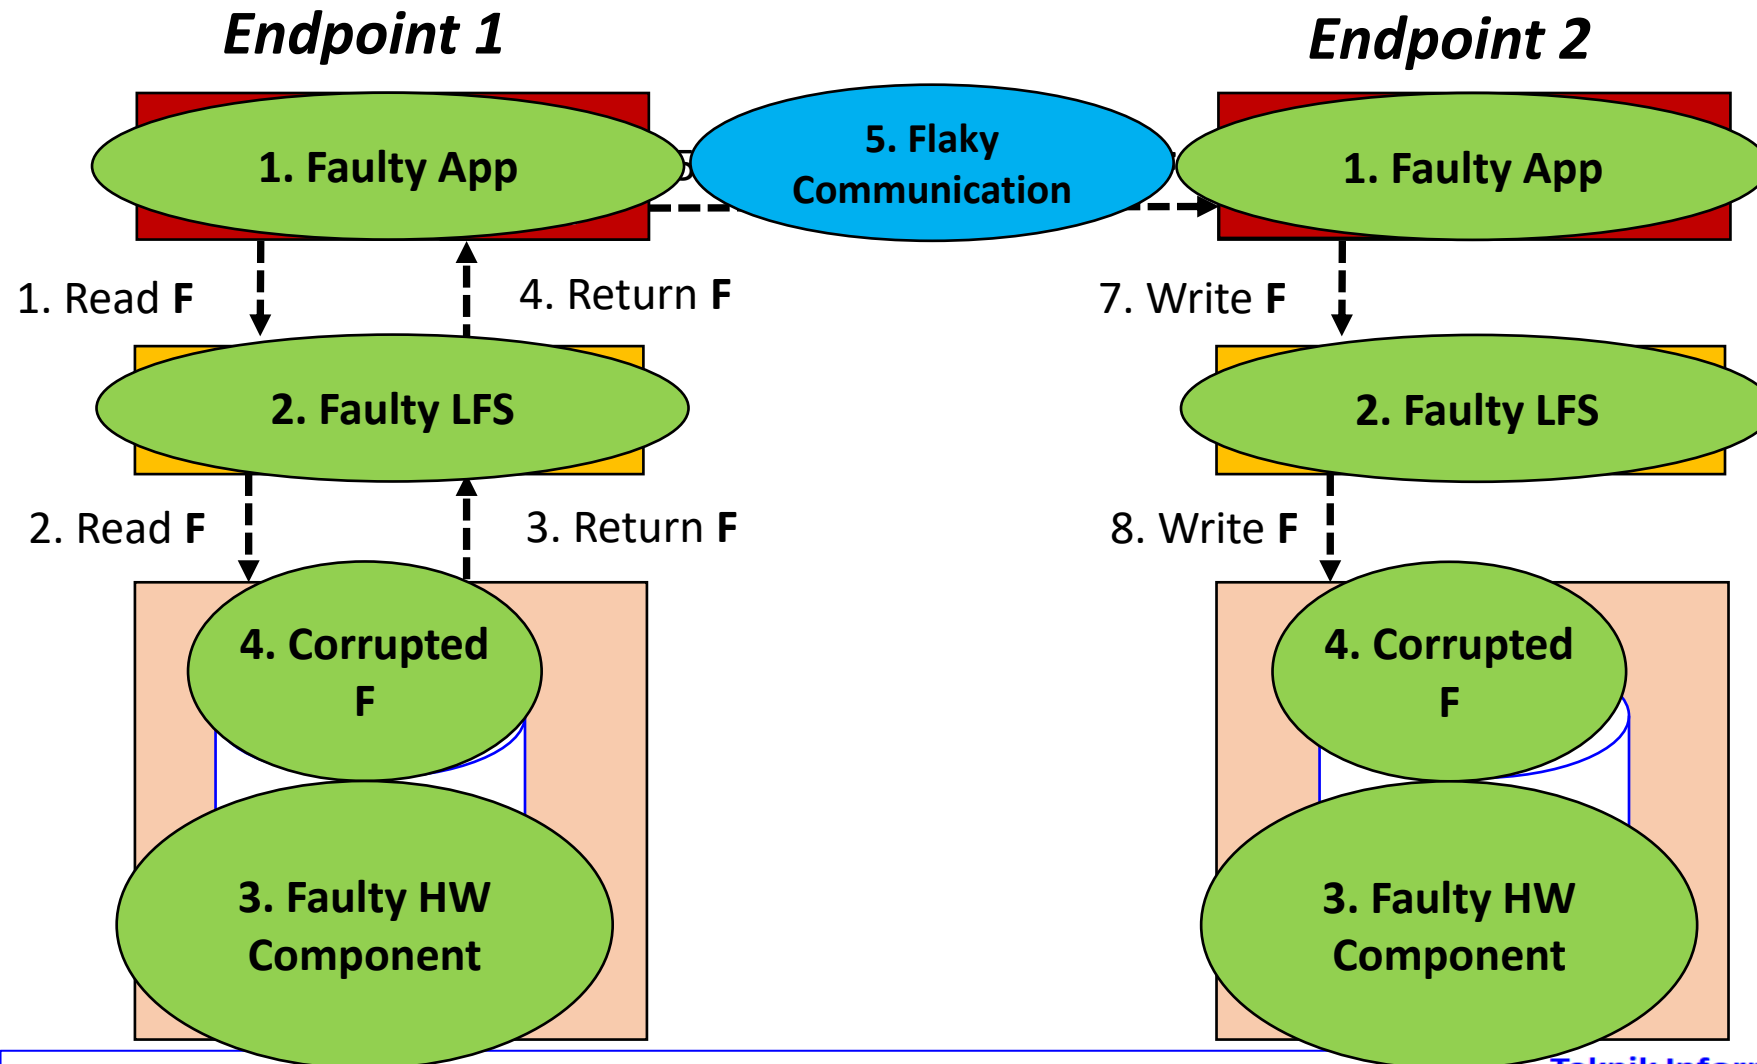
- If RPC is layered on top of UDP
 - Retransmission shall/can be handled by RPC
- If RPC is layered on top of TCP
 - Retransmission will be handled by TCP
 - Is it still necessary to take fault-tolerance measures within RPC?
 - Yes-- read “End-to-End Arguments in System Design” by Saltzer *et. al.*

Careful File Transfer: Flow



DCS = Data Communication System; LFS = Local File System

Careful File Transfer: Possible Threats



DCS = Data Communication System; LFS = Local File System

Careful File Transfer: End-To-End Check and Retry

- Endpoint 1 stores with F a checksum C_A
- After Endpoint 2 writes F , it reads it again from disk, calculates a checksum C_B , and sends it back to Endpoint 1
- Endpoint 1 compares C_A and C_B
 - If $C_A = C_B$, commit the file transfer
 - Else, retry the file transfer

Careful File Transfer: End-To-End Check and Retry

- How many retries?
 - Usually 1 if failures are rare
 - 3 retries might indicate that some part of the system needs repair
- What if the Data Communication System uses TCP?
 - Only threat 5 (e.g., packet loss due to a flaky communication) is eliminated
 - The frequency of retries gets reduced if the fault was caused by the communication system
 - More *control* traffic, but only missing parts of F need to be reshipped
 - The file transfer application still needs to apply *end-to-end reliability measures!*

Careful File Transfer: End-To-End Check and Retry

- What if the Data Communication System uses UDP?
 - Threat 5 (e.g., packet loss due to a flaky communication) is NOT eliminated- *F needs to be reshipped by the application if no measures are taken to address this threat*
 - The frequency of retries might increase
 - Worse performance on flaky links
 - *The file transfer application still needs to apply end-to-end reliability measures!*

In *both cases*, the application needs to provide end-to-end reliability guarantees!

Kuliah Berikutnya

- Layanan Penamaan

Pertanyaan?